

Write up Project 3


December 11, 2023

```
[1]: from numpy import *
      from matplotlib.pyplot import *
      import matplotlib.pyplot as plt
      import util, dr, datasets, softmax, runClassifier
      from utils import *
      from softmax import *
      from runClassifier import *
```


Part 1: PCA

Qpca1: Implement PCA

```
[9]: Si = util.sqrt(array([[3,2],[22,4]]))
      x = dot(random.randn(1000,2), Si)
      plot(x[:,0], x[:,1], 'b.')
      show(False)
      dot(x.T,x)/real(x.shape[0] - 1)
```



```
[9]: array([[26.18828362, 12.38687024],  
          [12.38687024,  6.24247454]])
```





```
[10]: (P, Z, evals) = dr.pca(x,2)
      Z
```

```
[10]: array([[ 0.90201048, -0.43171413],
             [ 0.43171413,  0.90201048]])
```

```
[11]: evals
```

```
[11]: array([31.83153594,  0.31273565])
```

```
[12]: x0 = dot(dot(x, Z[:,0]).reshape(1000,1), Z[:,0].reshape(1,2))
      x1 = dot(dot(x, Z[:,1]).reshape(1000,1), Z[:,1].reshape(1,2))
      plot(x[:,0], x[:,1], 'b.', x0[:,0], x0[:,1], 'r.', x1[:,0], x1[:,1], 'g.')
      show(False)
```



```
[14]: (X,Y) = datasets.loadDigits()  
(P1,Z1,evals1) = dr.pca(X, 784)  
evals
```

```
[14]: array([ 5.47145902e-02,  4.32457401e-02,  3.91832356e-02,  3.07589801e-02,  
            2.97240699e-02,  2.48634922e-02,  1.98552656e-02,  1.75982431e-02,  
            1.57833269e-02,  1.31208707e-02,  1.27918802e-02,  1.18266157e-02,  
            1.12559221e-02,  1.00752373e-02,  9.36845932e-03,  8.73601732e-03,  
            8.27005062e-03,  7.92899274e-03,  7.23521657e-03,  6.85924774e-03,  
            6.42706546e-03,  6.14996142e-03,  5.91821028e-03,  5.67737906e-03,  
            5.40356132e-03,  5.30190215e-03,  5.02203749e-03,  4.59637089e-03,  
            4.42538322e-03,  4.36981194e-03,  4.16187116e-03,  4.11031674e-03,  
            3.77459052e-03,  3.62187854e-03,  3.41085954e-03,  3.28497607e-03,  
            3.23280900e-03,  3.19587689e-03,  3.02982002e-03,  2.87978872e-03,  
            2.83078667e-03,  2.74927351e-03,  2.62589089e-03,  2.54642734e-03,  
            2.47927464e-03,  2.37555612e-03,  2.29383593e-03,  2.28594556e-03,  
            2.20977922e-03,  2.08730066e-03,  1.99999581e-03,  1.96943065e-03,  
            1.92821820e-03,  1.88309395e-03,  1.84886396e-03,  1.76088832e-03,  
            1.67020455e-03,  1.61444912e-03,  1.59805234e-03,  1.57744026e-03,  
            1.48050938e-03,  1.44561925e-03,  1.43009511e-03,  1.40344365e-03,  
            1.37760752e-03,  1.34557952e-03,  1.31369874e-03,  1.29776443e-03,  
            1.26802185e-03,  1.20748247e-03,  1.19311640e-03,  1.13277494e-03,  
            1.10668220e-03,  1.08849526e-03,  1.06129719e-03,  1.04324556e-03,  
            1.01409526e-03,  1.00687447e-03,  9.81023790e-04,  9.38761970e-04,  
            9.14566689e-04,  9.11766083e-04,  8.98271843e-04,  8.61728296e-04,  
            8.51322737e-04,  8.37576333e-04,  8.07377310e-04,  7.89436310e-04,  
            7.77157242e-04,  7.65601500e-04,  7.43889898e-04,  7.36837174e-04,  
            7.23782179e-04,  7.16183052e-04,  6.88886751e-04,  6.76881361e-04,  
            6.73658879e-04,  6.57964557e-04,  6.31789458e-04,  6.14600872e-04,  
            6.06852214e-04,  5.97638691e-04,  5.82295974e-04,  5.75874171e-04,  
            5.69167770e-04,  5.66582916e-04,  5.56170605e-04,  5.46447592e-04,  
            5.35887081e-04,  5.29480645e-04,  5.12646375e-04,  5.10076375e-04,  
            4.98901433e-04,  4.92492964e-04,  4.83279552e-04,  4.77823523e-04,  
            4.74356288e-04,  4.58393177e-04,  4.47510129e-04,  4.41138118e-04,  
            4.34427000e-04,  4.26534001e-04,  4.22078706e-04,  4.20811833e-04,  
            4.16611215e-04,  4.07636898e-04,  4.01008493e-04,  3.88768112e-04,  
            3.84665028e-04,  3.82008801e-04,  3.70120096e-04,  3.68936048e-04,  
            3.61495769e-04,  3.52564002e-04,  3.50413432e-04,  3.43551550e-04,  
            3.42456916e-04,  3.34070452e-04,  3.31285706e-04,  3.25466266e-04,  
            3.23319301e-04,  3.21575937e-04,  3.13465865e-04,  3.09027164e-04,  
            3.02370057e-04,  2.97280504e-04,  2.94665985e-04,  2.92498807e-04,  
            2.91135494e-04,  2.87403980e-04,  2.83721082e-04,  2.77237998e-04,  
            2.73678546e-04,  2.70133830e-04,  2.69075520e-04,  2.65467924e-04,  
            2.60293374e-04,  2.58426988e-04,  2.54266891e-04,  2.53192164e-04,
```

2.49764944e-04,	2.42950419e-04,	2.42151942e-04,	2.41786314e-04,
2.40184890e-04,	2.34384637e-04,	2.29678760e-04,	2.25975527e-04,
2.24052785e-04,	2.22648492e-04,	2.21183849e-04,	2.18103133e-04,
2.16772283e-04,	2.11643742e-04,	2.10557240e-04,	2.06598941e-04,
2.03725106e-04,	2.03002825e-04,	2.01737235e-04,	1.99838248e-04,
1.97020900e-04,	1.95882771e-04,	1.90553874e-04,	1.90035654e-04,
1.86202179e-04,	1.84195282e-04,	1.82248919e-04,	1.81374403e-04,
1.78634663e-04,	1.77003084e-04,	1.76734926e-04,	1.74827021e-04,
1.73861570e-04,	1.69074913e-04,	1.68301018e-04,	1.65926793e-04,
1.64763869e-04,	1.63508376e-04,	1.60460373e-04,	1.58903225e-04,
1.57550898e-04,	1.57069298e-04,	1.54714752e-04,	1.53382367e-04,
1.52196112e-04,	1.51851924e-04,	1.49036629e-04,	1.47581845e-04,
1.47015964e-04,	1.45691497e-04,	1.44753088e-04,	1.41471392e-04,
1.39746586e-04,	1.39175790e-04,	1.37987629e-04,	1.36408508e-04,
1.33085110e-04,	1.32389884e-04,	1.32027803e-04,	1.29974745e-04,
1.28716136e-04,	1.28319783e-04,	1.27568188e-04,	1.25324466e-04,
1.24347102e-04,	1.22744275e-04,	1.22300425e-04,	1.20760186e-04,
1.19779984e-04,	1.17751428e-04,	1.16606506e-04,	1.15922466e-04,
1.15017447e-04,	1.12928827e-04,	1.11734652e-04,	1.11213415e-04,
1.10236654e-04,	1.08918378e-04,	1.08252462e-04,	1.06430084e-04,
1.04553397e-04,	1.04281598e-04,	1.04054083e-04,	1.01576056e-04,
1.01343340e-04,	1.01267465e-04,	1.00522423e-04,	9.85053164e-05,
9.78816815e-05,	9.70690379e-05,	9.62925456e-05,	9.55812780e-05,
9.45559635e-05,	9.36019509e-05,	9.24278471e-05,	9.19656175e-05,
9.14391679e-05,	8.96372447e-05,	8.93024179e-05,	8.75651715e-05,
8.69322559e-05,	8.54624033e-05,	8.52465643e-05,	8.45712126e-05,
8.42354339e-05,	8.35126534e-05,	8.26637951e-05,	8.16250614e-05,
8.10797669e-05,	7.96140119e-05,	7.94147014e-05,	7.88650653e-05,
7.76851426e-05,	7.69168134e-05,	7.59915224e-05,	7.54955439e-05,
7.46249502e-05,	7.41006358e-05,	7.36929743e-05,	7.23523711e-05,
7.20429520e-05,	7.07902291e-05,	7.06927386e-05,	7.01269353e-05,
6.96346942e-05,	6.87641980e-05,	6.82166674e-05,	6.76223162e-05,
6.70368715e-05,	6.65200035e-05,	6.56722876e-05,	6.51808836e-05,
6.42906211e-05,	6.36604944e-05,	6.29414925e-05,	6.27787655e-05,
6.12849305e-05,	6.09263075e-05,	6.01884006e-05,	5.95464554e-05,
5.80776976e-05,	5.78205279e-05,	5.75292392e-05,	5.69274446e-05,
5.65771276e-05,	5.64412985e-05,	5.57036459e-05,	5.50499316e-05,
5.49942491e-05,	5.36323282e-05,	5.30991777e-05,	5.26568756e-05,
5.23712794e-05,	5.19583484e-05,	5.18381924e-05,	5.15401960e-05,
5.04554852e-05,	4.90569702e-05,	4.86487355e-05,	4.83589125e-05,
4.76003331e-05,	4.70510781e-05,	4.59385460e-05,	4.57307845e-05,
4.52785302e-05,	4.48145357e-05,	4.46440287e-05,	4.40717511e-05,
4.36747605e-05,	4.32923298e-05,	4.29148619e-05,	4.26332825e-05,
4.19585597e-05,	4.16284803e-05,	4.11736824e-05,	4.06097570e-05,
4.04030818e-05,	4.01164886e-05,	3.95221559e-05,	3.92048548e-05,
3.88512263e-05,	3.85824261e-05,	3.82123548e-05,	3.75291417e-05,
3.72292299e-05,	3.69259350e-05,	3.65790506e-05,	3.62293599e-05,

3.61688812e-05,	3.57178110e-05,	3.52504156e-05,	3.46303776e-05,
3.43365755e-05,	3.37485856e-05,	3.36655102e-05,	3.30809444e-05,
3.29371687e-05,	3.22457276e-05,	3.17451188e-05,	3.11110595e-05,
3.10819608e-05,	3.06394735e-05,	3.01036807e-05,	2.99597482e-05,
2.94869075e-05,	2.92286816e-05,	2.89553410e-05,	2.86577988e-05,
2.81775649e-05,	2.78537419e-05,	2.76491103e-05,	2.74490176e-05,
2.69949143e-05,	2.66711158e-05,	2.63043530e-05,	2.59990966e-05,
2.55178324e-05,	2.54420060e-05,	2.49967869e-05,	2.47386118e-05,
2.44478402e-05,	2.42861628e-05,	2.37778737e-05,	2.34678050e-05,
2.29924394e-05,	2.27444158e-05,	2.25710993e-05,	2.21925530e-05,
2.16898469e-05,	2.12574642e-05,	2.10710698e-05,	2.06125368e-05,
2.04801867e-05,	2.03528951e-05,	2.01852754e-05,	2.00109236e-05,
1.96358067e-05,	1.95357198e-05,	1.88062148e-05,	1.84740717e-05,
1.82343730e-05,	1.79672046e-05,	1.78883200e-05,	1.77158139e-05,
1.75571151e-05,	1.72617510e-05,	1.70780959e-05,	1.67246044e-05,
1.65874069e-05,	1.64043907e-05,	1.61572993e-05,	1.59549823e-05,
1.56642274e-05,	1.55508890e-05,	1.52751700e-05,	1.50113639e-05,
1.45628889e-05,	1.43395732e-05,	1.40217384e-05,	1.37922572e-05,
1.37149459e-05,	1.34724106e-05,	1.34004024e-05,	1.31395277e-05,
1.28611756e-05,	1.27115880e-05,	1.25654559e-05,	1.23064597e-05,
1.21871093e-05,	1.20404242e-05,	1.19082819e-05,	1.17445668e-05,
1.14430627e-05,	1.13639877e-05,	1.09683346e-05,	1.07615238e-05,
1.06384075e-05,	1.05582384e-05,	1.03033757e-05,	1.01843376e-05,
9.99360386e-06,	9.88765753e-06,	9.45538175e-06,	9.43379699e-06,
9.33984633e-06,	9.25979331e-06,	9.00558539e-06,	8.88594799e-06,
8.76931369e-06,	8.67576620e-06,	8.44589675e-06,	8.42047973e-06,
8.29299911e-06,	8.18500338e-06,	8.10426232e-06,	7.86968087e-06,
7.80326767e-06,	7.61047354e-06,	7.54069273e-06,	7.32839510e-06,
7.02790061e-06,	6.93803992e-06,	6.86839425e-06,	6.81302794e-06,
6.72974568e-06,	6.61403985e-06,	6.50509538e-06,	6.33303381e-06,
6.22118912e-06,	5.97738030e-06,	5.81367159e-06,	5.66976470e-06,
5.57648468e-06,	5.54174989e-06,	5.42658176e-06,	5.19783419e-06,
5.13442795e-06,	5.06890822e-06,	4.98959410e-06,	4.82534928e-06,
4.66035688e-06,	4.56198566e-06,	4.51001805e-06,	4.36820868e-06,
4.29161273e-06,	4.24292681e-06,	4.13953158e-06,	4.08394845e-06,
3.84565097e-06,	3.81561572e-06,	3.68911200e-06,	3.65366772e-06,
3.62286115e-06,	3.48300688e-06,	3.41704003e-06,	3.26430511e-06,
3.21058430e-06,	3.14052045e-06,	3.06495818e-06,	3.01867376e-06,
2.89121200e-06,	2.82203991e-06,	2.76596354e-06,	2.68021341e-06,
2.57159092e-06,	2.48812905e-06,	2.43153349e-06,	2.37842654e-06,
2.26778695e-06,	2.20089413e-06,	2.12019860e-06,	2.03927844e-06,
2.00474900e-06,	1.96567566e-06,	1.86980680e-06,	1.85320127e-06,
1.80781363e-06,	1.75081389e-06,	1.69244461e-06,	1.64184031e-06,
1.60386521e-06,	1.57252147e-06,	1.49681811e-06,	1.39316172e-06,
1.32824152e-06,	1.27965819e-06,	1.22198261e-06,	1.16411084e-06,
1.14888666e-06,	1.07544038e-06,	1.03222807e-06,	9.91356743e-07,
9.64565963e-07,	9.55180458e-07,	8.87585979e-07,	8.67379177e-07,

[illegible]


```

0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
0.00000000e+00, 0.00000000e+00, -1.30009069e-22, -4.78099872e-21,
-4.78099872e-21, -5.03227736e-21, -5.03227736e-21, -8.00168618e-21,
-1.39413532e-20, -1.39413532e-20, -3.56875226e-20, -3.56875226e-20,
-4.05460168e-20, -4.87670768e-20, -4.87670768e-20, -7.37491520e-20])

```

Qcpa2


```

[18]: #code for Qpca2:
normal_evals = evals1/sum(evals)
index = range(0, len(normal_evals))
plot(index, normal_evals)
xlabel('index')
ylabel('eigenvalues')
title('Normalized eigenvalues')

cum_sum = cumsum(normal_evals)
print(argmax(cum_sum > 0.9))
print(argmax(cum_sum > 0.95))



```

81
135



There are 81 eigenvectors we need to include before we've accounted for 90% of the variance. For 95% of the variance, the number goes up to 135.

```
[17]: #Qpca3:  
util.drawDigits(Z.T[:50,:], arange(50))  
show(False)
```



Qpca3: These graphs do not exactly look like digit. Some of them somewhat resemble the digits but most of them are not very clear. They are being blurry due to the fact that projecting the eigenvectors onto the top 50 eigenvectors yields 82.7% accuracy of the total variance of the dataset. Hence, we compress the image to a lower dimension, which explains why the graphs are not clearly defined.

Part 2: Softmax Regression.

Qsr1:

1) Given the probability: $P = [y = i] = \frac{e^{\vec{w}_i \cdot \vec{x}}}{\sum_j e^{\vec{w}_j \cdot \vec{x}}}$

Sum of the probabilities:

$$\sum_i P = [y = i] = \sum_i \frac{e^{\vec{w}_i \cdot \vec{x}}}{\sum_j e^{\vec{w}_j \cdot \vec{x}}} = \frac{\sum_i e^{\vec{w}_i \cdot \vec{x}}}{\sum_j e^{\vec{w}_j \cdot \vec{x}}} = 1$$

Hence, the probability sum to 1.

2) The dimension of W is M x C, where C is the number of possible classification and M is the number of input features. The dimension of X is N x M where N is number of examples. The dimension of WX is C x N.

Qsr3:

1) We have: $P = [y = i] = \frac{e^{\vec{w}_i \cdot \vec{x}}}{\sum_j e^{\vec{w}_j \cdot \vec{x}}}$

Now, let $\max(\vec{w}_i x) = WX_{max}$. We divide the $e^{WX_{max}}$ to the probability, result:

$$P = [y = i] = \frac{\frac{e^{\vec{w}_i \cdot \vec{x}}}{e^{WX_{max}}}}{\frac{\sum_j e^{\vec{w}_j \cdot \vec{x}}}{e^{WX_{max}}}} = \frac{e^{\vec{w}_i \cdot \vec{x} - WX_{max}}}{\sum_j e^{\vec{w}_j \cdot \vec{x} - WX_{max}}}$$

This is still the same ratio, which means it does not affected the predicted accuracy.

2) This might be an optimization over using W_X because when we calculate the cost, a large value x can create a wrong value for total cost. The exponent of a large number is very significant. Hence, when divide with the exponent of a large number, the calculation might go wrong. Therefore, using this function might be better.

```
[4]: #code for Qrs4:
# MNIST images are 28 * 28
exSize = 28*28
# 10 digits
numClasses = 10
# Regularizer coefficient
X, Y = loadMNIST('data/train-images.idx3-ubyte', 'data/train-labels.idx1-ubyte')
sm = SoftmaxRegression(numClasses, exSize)
testX, testY = loadMNIST('data/t10k-images.idx3-ubyte', 'data/t10k-labels.
↳idx1-ubyte')
predictions = sm.predict(X)
(data_size, train_acc, test_acc) = learningCurve(sm, numClasses, exSize, X, Y,
↳testX, testY)
print(data_size)
```

60000

16

1

(784, 2)

(2,)

Training classifier on 2 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 3.45388D+00 |proj g|= 4.50000D-01

At iterate 1 f= 7.54791D-01 |proj g|= 2.56845D-01

At iterate 2 f= 7.35419D-01 |proj g|= 2.49947D-01

At iterate 3 f= 6.93397D-01 |proj g|= 2.49922D-01

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	4	59	2	0	0	2.499D-01	6.934D-01
F = 0.69339723797348007							

ABNORMAL_TERMINATION_IN_LNSRCH

Training accuracy 1, test accuracy 0.1139

2

(784, 4)

(4,)

Training classifier on 4 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

```

At iterate    0    f=  1.40360D+00    |proj g|=  4.37973D-01
At iterate    1    f=  3.61703D-01    |proj g|=  1.24020D-01
At iterate    2    f=  3.46576D-01    |proj g|=  1.24999D-01

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	3	51	2	0	0	1.250D-01	3.466D-01

F = 0.34657562925118829

```

ABNORMAL_TERMINATION_IN_LNSRCH
Training accuracy 1, test accuracy 0.1437
3
(784, 8)
(8,)
Training classifier on 8 points...
RUNNING THE L-BFGS-B CODE

```

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

```

At iterate    0    f=  1.77493D+01    |proj g|=  2.09057D-01
At iterate    1    f=  1.53034D+01    |proj g|=  1.95530D-01
At iterate    2    f=  2.03323D+00    |proj g|=  1.25000D-01
At iterate    3    f=  3.48882D-01    |proj g|=  1.16768D-01

```

...

```

At iterate   12    f=  2.39576D-05    |proj g|=  2.38353D-05

```

At iterate 13 f= 1.20328D-05 |proj g|= 1.19597D-05

At iterate 14 f= 6.05174D-06 |proj g|= 6.00330D-06

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	14	15	1	0	0	6.003D-06	6.052D-06
F = 6.0517408407289446E-006							

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.2692

4

(784, 15)

(15,)

Training classifier on 15 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 5.11229D+00 |proj g|= 1.76880D-01

At iterate 1 f= 3.35493D+00 |proj g|= 1.26666D-01

At iterate 2 f= 2.27328D+00 |proj g|= 1.55596D-01

At iterate 3 f= 1.24779D-01 |proj g|= 6.98760D-02

...

At iterate 12 f= 2.29767D-05 |proj g|= 1.01193D-05

At iterate 13 f= 1.22311D-05 |proj g|= 5.25885D-06

* * *

Tit = total number of iterations
Tnf = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip = number of BFGS updates skipped
Nact = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	13	14	1	0	0	5.259D-06	1.223D-05
F =	1.2231068946561243E-005						

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.3502

5

(784, 30)

(30,)

Training classifier on 30 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 5.69010D+00 |proj g|= 9.86926D-02

At iterate 1 f= 4.51942D+00 |proj g|= 9.86915D-02

At iterate 2 f= 1.44265D+00 |proj g|= 1.19933D-01

At iterate 3 f= 2.35187D-01 |proj g|= 5.25363D-02

At iterate 4 f= 1.11300D-01 |proj g|= 3.17428D-02

...

At iterate 12 f= 1.24037D-04 |proj g|= 2.94642D-05

At iterate 13 f= 6.69782D-05 |proj g|= 1.52075D-05

At iterate 14 f= 3.69401D-05 |proj g|= 8.24258D-06

* * *

Tit = total number of iterations
Tnf = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip = number of BFGS updates skipped
Nact = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	14	15	1	0	0	8.243D-06	3.694D-05
F =	3.6940080143739086E-005						

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.4004

6

(784, 59)

(59,)

Training classifier on 59 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.63305D+00 |proj g|= 8.05740D-02

At iterate 1 f= 1.77657D+00 |proj g|= 4.96180D-02

At iterate 2 f= 1.08953D+00 |proj g|= 6.26937D-02

At iterate 3 f= 6.01203D-01 |proj g|= 2.53500D-02

...

At iterate 13 f= 2.41959D-04 |proj g|= 3.23871D-05

At iterate 14 f= 1.29249D-04 |proj g|= 1.74530D-05

At iterate 15 f= 8.02618D-05 |proj g|= 3.50358D-05

At iterate 16 f= 3.05624D-05 |proj g|= 4.88372D-06

* * *

Tit = total number of iterations
Tnf = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip = number of BFGS updates skipped
Nact = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	16	17	1	0	0	4.884D-06	3.056D-05
F =	3.0562412316286158E-005						

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.5819

7

(784, 118)

(118,)

Training classifier on 118 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.99178D+00 |proj g|= 4.68696D-02

At iterate 1 f= 1.44924D+00 |proj g|= 3.09946D-02

At iterate 2 f= 9.85280D-01 |proj g|= 6.24452D-02

At iterate 3 f= 6.45540D-01 |proj g|= 3.73655D-02

At iterate 4 f= 4.11109D-01 |proj g|= 2.28940D-02

...

At iterate 19 f= 5.57072D-05 |proj g|= 1.95780D-05

At iterate 20 f= 2.69800D-05 |proj g|= 4.44389D-06

* * *

Tit = total number of iterations
Tnf = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip = number of BFGS updates skipped
Nact = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	20	21	1	0	0	4.444D-06	2.698D-05
F = 2.6980015531539907E-005							

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.7018

8

(784, 235)

(235,)

Training classifier on 235 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.47343D+00 |proj g|= 3.22743D-02

At iterate 1 f= 1.18133D+00 |proj g|= 2.45324D-02

At iterate 2 f= 8.43887D-01 |proj g|= 2.67739D-02

At iterate 3 f= 5.21751D-01 |proj g|= 4.18095D-02

At iterate 4 f= 2.89418D-01 |proj g|= 1.23715D-02

At iterate 5 f= 2.26249D-01 |proj g|= 8.73027D-03

...

At iterate 20 f= 1.33032D-04 |proj g|= 1.12012D-05

At iterate 21 f= 9.21129D-05 |proj g|= 2.01157D-05

At iterate 22 f= 5.19530D-05 |proj g|= 5.48372D-06

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	22	23	1	0	0	5.484D-06	5.195D-05
F =	5.1952951981449447E-005						

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.793

9

(784, 469)

(469,)

Training classifier on 469 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.04250D+00 |proj g|= 2.10990D-02

At iterate 1 f= 8.31139D-01 |proj g|= 1.53740D-02

At iterate 2 f= 6.07346D-01 |proj g|= 2.75524D-02

At iterate 3 f= 4.19349D-01 |proj g|= 1.48126D-02

At iterate 4 f= 2.75766D-01 |proj g|= 1.25437D-02

...

At iterate 27 f= 1.15651D-04 |proj g|= 1.31001D-05

At iterate 28 f= 9.89895D-05 |proj g|= 1.00974D-05

At iterate 29 f= 7.66305D-05 |proj g|= 6.20370D-06

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	29	30	1	0	0	6.204D-06	7.663D-05
F = 7.6630544212003617E-005							

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.815

10

(784, 938)

(938,)

Training classifier on 938 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 8.84480D-01 |proj g|= 8.93479D-03

At iterate 1 f= 7.71158D-01 |proj g|= 7.56693D-03

At iterate 2 f= 6.27666D-01 |proj g|= 3.31254D-02

...

At iterate 36 f= 1.73244D-04 |proj g|= 3.16258D-05

At iterate 37 f= 1.24814D-04 |proj g|= 1.19187D-05

At iterate 38 f= 1.10069D-04 |proj g|= 7.07214D-06

* * *

Tit = total number of iterations
Tnf = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip = number of BFGS updates skipped
Nact = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	38	39	1	0	0	7.072D-06	1.101D-04
F =	1.1006897943093058E-004						

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.8428

11

(784, 1875)

(1875,)

Training classifier on 1875 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 7.50555D-01 |proj g|= 8.62377D-03

At iterate 1 f= 6.70055D-01 |proj g|= 4.48365D-03

At iterate 2 f= 6.25365D-01 |proj g|= 4.71615D-03

At iterate 3 f= 5.23282D-01 |proj g|= 1.02471D-02

...

At iterate 60 f= 1.88922D-04 |proj g|= 1.53696D-05

At iterate 61 f= 1.68409D-04 |proj g|= 1.06619D-05

At iterate 62 f= 1.35620D-04 |proj g|= 6.19396D-06

* * *

Tit = total number of iterations
Tnf = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip = number of BFGS updates skipped
Nact = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	62	66	1	0	0	6.194D-06	1.356D-04
F = 1.3562009626588901E-004							

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.859

12

(784, 3750)

(3750,)

Training classifier on 3750 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 9.47158D-01 |proj g|= 8.23798D-03

At iterate 1 f= 9.05725D-01 |proj g|= 5.58973D-03

At iterate 2 f= 8.71385D-01 |proj g|= 5.17801D-03

...

At iterate 7 f= 5.51246D-01 |proj g|= 3.91548D-03

At iterate 8 f= 5.04666D-01 |proj g|= 8.69213D-03

This problem is unconstrained.

At iterate 9 f= 4.65868D-01 |proj g|= 8.26349D-03

At iterate 10 f= 4.44192D-01 |proj g|= 3.42502D-03

...

At iterate 133 f= 1.78186D-04 |proj g|= 1.13796D-05

At iterate 134 f= 1.65219D-04 |proj g|= 7.48054D-06

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	134	144	1	0	0	7.481D-06	1.652D-04

F = 1.6521861880642840E-004

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.8471

13

(784, 7500)

(7500,)

Training classifier on 7500 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

This problem is unconstrained.

At iterate 0 f= 1.23260D+00 |proj g|= 4.40204D-03

At iterate 1 f= 1.19851D+00 |proj g|= 4.23543D-03

At iterate 2 f= 1.17325D+00 |proj g|= 3.22072D-03

At iterate 3 f= 1.08138D+00 |proj g|= 6.82248D-03

At iterate 4 f= 1.01066D+00 |proj g|= 4.20362D-03

...

At iterate 341 f= 2.99889D-04 |proj g|= 1.24309D-05

At iterate 342 f= 2.82909D-04 |proj g|= 9.70383D-06

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	342	362	1	0	0	9.704D-06	2.829D-04
F = 2.8290915858524717E-004							

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

Training accuracy 1, test accuracy 0.8501

14

(784, 15000)

(15000,)

Training classifier on 15000 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 1.83205D+00 |proj g|= 3.84989D-03

At iterate 1 f= 1.79612D+00 |proj g|= 2.49122D-03

...

At iterate 7 f= 1.54689D+00 |proj g|= 1.98857D-03

At iterate 8 f= 1.51754D+00 |proj g|= 3.01389D-03

At iterate 9 f= 1.47373D+00 |proj g|= 4.94200D-03

This problem is unconstrained.

At iterate 10 f= 1.42990D+00 |proj g|= 3.76280D-03

At iterate 11 f= 1.40955D+00 |proj g|= 2.11316D-03

...

At iterate 398 f= 9.84071D-02 |proj g|= 2.75509D-04

At iterate 399 f= 9.81845D-02 |proj g|= 4.41805D-04

At iterate 400 f= 9.79706D-02 |proj g|= 3.73465D-04

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	400	427	1	0	0	3.735D-04	9.797D-02
F =	9.7970595972840205E-002						

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT

Training accuracy 0.9766, test accuracy 0.861

15

(784, 30000)

(30000,)

Training classifier on 30000 points...

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

```

At iterate    0    f=  1.00383D+00    |proj g|=  3.19584D-03
At iterate    1    f=  1.00016D+00    |proj g|=  2.96458D-03
At iterate    2    f=  9.96332D-01    |proj g|=  2.63122D-03
At iterate    3    f=  9.77399D-01    |proj g|=  3.05127D-03
At iterate    4    f=  9.70435D-01    |proj g|=  4.43586D-03

```

This problem is unconstrained.

```

At iterate    5    f=  9.61265D-01    |proj g|=  2.73159D-03
At iterate    6    f=  9.50354D-01    |proj g|=  1.62143D-03
...
At iterate  398    f=  2.24659D-01    |proj g|=  1.30028D-03
At iterate  399    f=  2.24426D-01    |proj g|=  6.81391D-04
At iterate  400    f=  2.24219D-01    |proj g|=  3.18458D-04

```

* * *

```

Tit   = total number of iterations
Tnf   = total number of function evaluations
Tnint = total number of segments explored during Cauchy searches
Skip  = number of BFGS updates skipped
Nact  = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F     = final function value

```

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	400	434	1	0	0	3.185D-04	2.242D-01

F = 0.22421920094254763

```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT
Training accuracy 0.9388, test accuracy 0.8832
16
(784, 60000)
(60000,)
Training classifier on 60000 points...
RUNNING THE L-BFGS-B CODE

```

* * *

Machine precision = 2.220D-16

N = 7840 M = 10

At X0 0 variables are exactly at the bounds

At iterate 0 f= 5.63363D-01 |proj g|= 7.75851D-03

At iterate 1 f= 5.43949D-01 |proj g|= 6.49524D-03

This problem is unconstrained.

At iterate 2 f= 5.31404D-01 |proj g|= 4.74214D-03

At iterate 3 f= 5.02349D-01 |proj g|= 3.67615D-03

At iterate 4 f= 4.98815D-01 |proj g|= 1.23197D-02

At iterate 5 f= 4.84652D-01 |proj g|= 4.27436D-03

At iterate 6 f= 4.78740D-01 |proj g|= 1.73815D-03

...

At iterate 398 f= 2.26960D-01 |proj g|= 4.39289D-04

At iterate 399 f= 2.26853D-01 |proj g|= 2.77452D-04

At iterate 400 f= 2.26781D-01 |proj g|= 3.58551D-04

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *


N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
7840	400	424	1	0	0	3.586D-04	2.268D-01
F = 0.22678116337494983							

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT

Training accuracy 0.93805, test accuracy 0.9172

```
[2.000e+00 4.000e+00 8.000e+00 1.500e+01 3.000e+01 5.900e+01 1.180e+02  
2.350e+02 4.690e+02 9.380e+02 1.875e+03 3.750e+03 7.500e+03 1.500e+04  
3.000e+04 6.000e+04]
```

```
[6]: plt.plot(train_acc, 'ro')  
plt.plot(test_acc, 'bs')  
plt.ylabel('Accuracy')  
plt.xlabel('Examples Seen')  
plt.legend(['Training Accuracy', 'Test Accuracy'])  
plt.title('Training and Test Accuracy vs. Examples Seen')  
plt.show()
```



QRS4: When the examples is low, the overfit occurs. When there are more examples, the overfit decreases. When it passes 14, the accuracy of training and testing are the same.

Part 3: NN

Qnn1.1: Submitted as code

Qnn1.2: Submitted as code

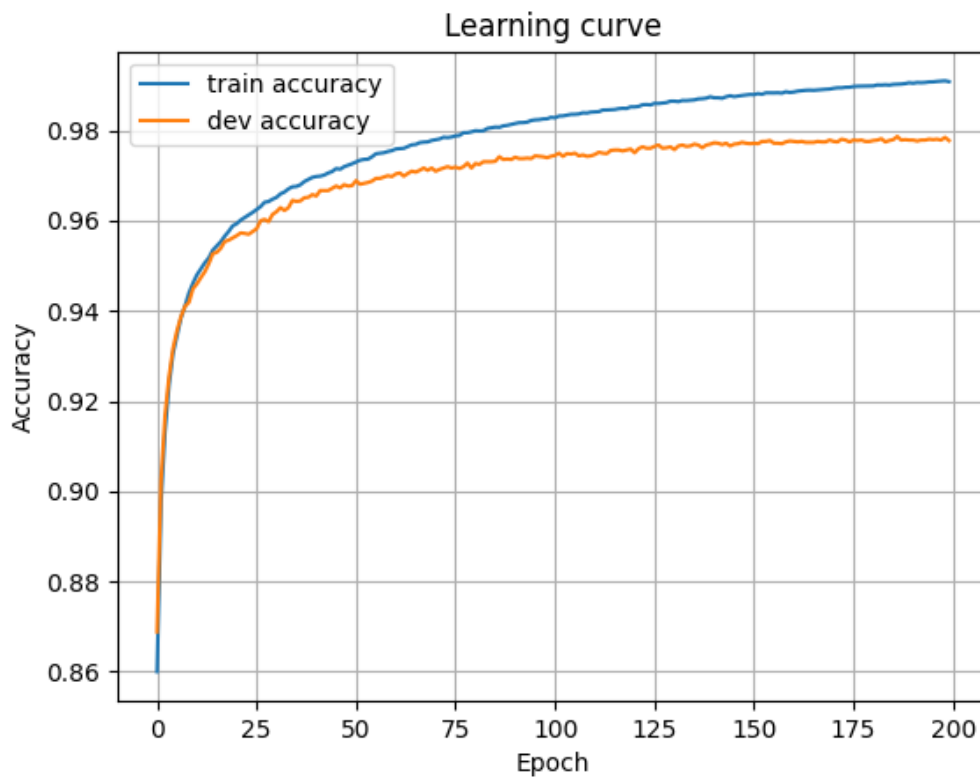
Qnn1.3:

With the following settings, the final dev accuracy for 200/200 epoch is **0.979**.

Settings:

1. Number of layers: 2
2. Size of layers 1 and 2: [196, 196]
3. Activation function: ReLU
4. Loss function: Squared Loss
5. Learning rate: 0.01
6. batch size: 64
7. Max epochs: 200

The Training and Dev accuracy for NN with the settings above is:



Qnn1.4:

In some cases it may be appropriate to initialize the entries of the weight matrix as small random numbers rather than all zeros because:

- **We need to avoid dead neurons.** If all weights are initialized to zero, the neurons in a layer will all produce the same output, and during backpropagation, they will all receive the same gradients. For example, we have the first 2 layers of multi-layer perceptron. During forward propagation, each unit in the hidden layer gets a signal $\sum_i^N W_{i,j} \cdot x_i$. It does not matter what input we get, if the weights are the same, all the hidden units will be the same. This can lead to a situation where neurons become "dead" (i.e., they always output the same values) and do not contribute meaningfully to the learning process. Therefore, random initialization helps prevent this by ensuring that neurons start with different values, encouraging them to learn distinct features.

Extra Credit

NOTE 1: We implemented the new optimizer, the **AMSGrad**, in **direction 3**. The modified implementation is in:

1. **nn_extra.py**: modified update rules
2. **run_nn_extra.py**: run experiments

NOTE 2:: In the implementation, we first run a comparative study between AMSGrad and SGD on different training set sizes (e.g: 1/100 or 1/2 original size) to test its convergence ability with different training data available. The dev set size is kept the same. It can take about 5 minutes to train 2 models on all different sizes.

Qnn2.1:

- What I did was re-implementing the **update()** function in the NN class by adding the 1st and 2nd order momentum and using them by using the following update rule (**AMSGrad()**):

$$\begin{aligned}
 m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\
 v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \\
 \hat{v}_t &= \max(\hat{v}_{t-1}, v_t) \\
 \theta_{t+1} &= \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} m_t
 \end{aligned}$$

- Our experiments show that **AMSGrad** requires less sample to converge compared to **SGD**. The experiments are shown below:

Optimizer Type	1/100 train set	1/10 train set	1/5 train set	1/2 train set	full training set
SGD	0.7561	0.9184	0.9507	0.9619	0.9725
AMSGrad	0.8592	0.94	0.9568	0.9686	0.9781

- Our experiments below show that AMSGrad has **significantly faster convergence rate** and **better overall performance** compared to SGD with the same architecture, loss function, and optimal hyperparameters. The comparative study on the convergence rate between AMSGrad and SGD (in terms of training loss, training accuracy, and dev accuracy) is shown in the 2 plots below:

